

Expert Grammaticality Judges Are Evaluators, Not Participants

Grammaticality Judgments, Rater Roles, and Human-Subjects Review

Brett Reynolds *

Humber Polytechnic & University of Toronto

19th June 2026

Abstract

Expert judgments of grammaticality are always human acts, but they are not always participant data. When a linguist is asked to evaluate linguistic materials in virtue of professional competence, the role is closer to that of an essay scorer, translation reviewer, coding adjudicator, or psychometric rater than to that of a participant whose own behaviour, attitudes, or cognitive processing are the target of study. This paper argues for a role-sensitive distinction between participant data and expert evaluation in grammaticality research. The distinction does not remove ordinary obligations of transparency, consent, conflict management, or local ethics review where required. It instead clarifies what the evidence is evidence of: not a sample of speaker reactions, but a disciplined expert assessment of linguistic materials under stated standards.

Keywords: grammaticality judgments; expert judgment; rater effects; linguistic methodology; research ethics

I THE PROBLEM

Ask whether the grammaticality judgments in a syntax paper needed consent and ethics review. The answer turns on what the question hides: whether the study is about that person, not just whether a person produced the judgment. Expert grammaticality judgments occupy an unstable methodological category. In one description, they are human responses to linguistic stimuli, and so look like participant data. In another, they are scholarly evaluations of linguistic materials, and so look like expert ratings, editorial judgments, coding decisions, or psychometric scoring.

The pull toward the first description has a respectable source. A strand of experimental syntax has argued that informal intuitions, often the author's own, are a weak evidential base, and that acceptability should be measured on samples of naive speakers under controlled conditions (Gibson & Fedorenko, 2010, 2013; Gibson et al., 2011). On that view the judges are a sample, their responses

*Contact: brett.reynolds@humber.ca

are the data, and the design is participant research. The worry is real, and the corrective has improved the field.

The corrective answers a measurement question, not a role question. That informal intuitions can be unreliable is a reason to calibrate, sample, or formalise the measurement, and the reliability of informal judgments has itself been tested and largely upheld (Schütze, 2016; Sprouse et al., 2013). It is not a reason to treat everyone who supplies a grammaticality judgment as the object of study. Reliability is a property of the instrument; it does not fix what the study is about.

So the answer is not that expert judgment escapes scrutiny, but that the scrutiny should follow the role. A judge is a participant when the question concerns that person's reaction, processing, background, identity, or distribution of responses. A judge is an evaluator when the question concerns the linguistic status of an item and expertise is used to assess it against disciplinary standards. Misclassifying costs both ways: treating evaluation as sampling overstates what the judgment shows, and treating sampling as evaluation strips persons of protections they are owed.

The target claim is narrow. Expert judges of grammaticality should not be treated as participants merely because their judgments are human acts. The first question is what role the judgment plays in the evidential design.

2 PARTICIPANTS, INFORMANTS, AND EVALUATORS

The central distinction is between producing data about oneself and applying expertise to a target object. A participant contributes evidence about a population of people, a psychological process, a distribution of acceptability, a demographic contrast, or an intervention effect. An evaluator contributes an assessment of an object: an essay, a translation, a diagnosis, a code, a performance, or a linguistic example.

The middle term, informant, names a recurring hard case rather than a third evidential role. A fieldwork consultant who supplies primary data about their own language is, by the test in the previous paragraph, contributing evidence about a speaker and a community, which places them with participants, not with essay scorers. The term turns slippery because the same person may later be asked to apply a stated criterion to a specified item, which is evaluation. Informant marks where the two roles meet, not a stable category between them.

This distinction is functional rather than honorific. Calling someone an evaluator does not make the judgment infallible, unregulated, or methodologically private. It changes the evidential question. The relevant issues become qualification, calibration, independence, adjudication, documentation, and conflict of interest, not sampling, participant burden, demographic representativeness, or the protection of private personal data.

What counts as EXPERTISE here is not a professional title alone but a role in a design. A judge is an evaluator when selected to apply a stated standard to a linguistic object, so the evidential weight of the judgment turns on qualification, independence, and task specification rather than on credentials.

The cases form a spectrum: authorial self-checking, informal consultation with a colleague, independent expert evaluation, and the systematic collection of judgments from a sample. A linguist reporting their own reaction as one of a population is not an evaluator in this sense; a linguist applying a disciplinary criterion to a specified item may be.

For grammaticality work, the distinction matters because many expert judgments do not estimate speaker-preference distributions. The task is to decide whether a proposed sentence, construction,

contrast, or diagnostic behaves as the theory says it does. That use of expertise is evaluative.

3 THE ESSAY-SCORER ANALOGY

Educational and psychometric measurement draws a line that grammaticality research often blurs: the scored are not the scorers. Essay scorers, oral-proficiency raters, and constructed-response markers can introduce error, bias, drift, and disagreement, and the field takes that seriously. None of it converts the scorers into participants alongside the students.

What makes the analogy bite is that this scoring is not casual. Constructed-response programmes train raters against rubrics, double-score responses, monitor severity and drift, adjudicate disagreements, and model rater effects statistically (American Educational Research Association et al., 2014; McCaffrey et al., 2022). The scoring is about as systematic, multi-rater, and quantitatively scrutinised as human judgment gets. The raters are still part of the measurement apparatus, not its object.

That ordering matters, because the usual worry about grammaticality judgments runs the other way. The intuition is that once judgment collection becomes systematic, with many judges, fixed materials, and reliability statistics, it must be participant research. The essay-scorer case shows the inference backwards. Systematicity, multiple raters, and reliability analysis are signatures of a measurement procedure, not marks of a study about the judges.

The analogy buys neither infallibility nor exemption from scrutiny. It relocates the scrutiny. If raters can be unreliable, the response is to train, blind, rotate, calibrate, adjudicate, report: the standards that govern scoring, not those that govern participant sampling (American Educational Research Association et al., 2014). An expert grammaticality judge deciding whether an item behaves as a theory predicts is doing the rater's job, and should be held to the rater's standards.

The analogy is not exact, and the gap cuts in a useful direction. Essay scorers normally judge performances produced independently of them, whereas grammaticality evaluators often judge examples constructed for the argument they support. That raises the need for disclosure, independence, and, where feasible, blinding to the prediction. It does not change the role; it changes the controls the role requires.

4 WHAT EXPERT GRAMMATICALITY JUDGMENTS DO

Expert grammaticality judgments can serve several roles. They can screen examples, test diagnostic contrasts, adjudicate corpus candidates, check paraphrases, evaluate minimal pairs, or assess whether a constructed item successfully isolates the intended variable. These are not the same role as estimating how a population would rate the sentence on a Likert scale.

The methodological burden is correspondingly different. An expert-judgment design should say what qualifies the judge, what the judge was asked to evaluate, what information was available, whether judgments were independent, how disagreements were handled, and whether the judgment bears on form, meaning, register, dialect, processing difficulty, or theoretical classification.

Level discipline matters here. A grammaticality claim, an acceptability-distribution claim, a processing claim, and a social-distribution claim are different claims. Treating every expert judgment as participant data blurs those levels instead of protecting them.

5 WHEN JUDGES ARE PARTICIPANTS

The evaluator role has boundaries. If the study asks how linguists respond, how expertise changes judgments, how quickly people process a contrast, how dialect background affects acceptability, or

how training changes ratings, then the judges belong to the target population. In such cases, the design studies their responses as responses, so they function as participants or participant-like contributors.

The clearest limiting case sits at the other end. A single-authored syntax article rests on the author's own judgments: the author builds the examples, marks them, and those judgments are the evidential core. If supplying a judgment relevant to the research question made one a participant, the author would be a participant in their own study, owing self-consent, a right to withdraw, and review of the author as a subject. No board asks this, and on that reading the field's pre-experimental history was unconsented self-experimentation. Producing the judgment cannot be the trigger.

Parity shows what the trigger is instead. Let the author judge fifty sentences and a colleague judge the same fifty; the judgments do identical evidential work, and only the producer differs. So the line falls on the focus of inquiry, not the human act. This does not make a lone intuition good evidence. That it can be unreliable is the measurement worry of [section 1](#), met by calibration and replication, not by recasting the author as a research subject.

The same person can also move between roles. A linguist who privately checks an example for a colleague may be an evaluator. The same linguist in an experiment comparing expert and non-expert ratings is a participant. The ethical and methodological classification follows the study role, not the person's professional title.

Because the role follows the design, it can look manipulable: relabel a judgment study as "evaluation" and the protections fall away. What blocks the move is that the role is fixed by the target of the research question, not by the label a researcher prefers. Take the hard case: thirty trained linguists rate two hundred sentences, and the reliability is reported. If the ratings certify the status of the items, the design is evaluation. If the distribution of ratings is itself the finding, the linguists are the object of study, and they are participants. The diagnostic is what the judgment is evidence of, settled before any label is applied.

The distinction cuts both ways, against overreach and against special pleading. Expertise does not exempt a study from human-participant review when the human is the object of inquiry. But neither does the mere presence of human expertise convert every scholarly evaluation into participant research.

6 WHAT THE REGULATIONS ACTUALLY SAY

Support for the distinction in the regulatory texts is uneven, and the gap is jurisdictional. United States policy builds the *from/about* contrast into its definition. Under the Common Rule, a human subject is a living individual "about whom" an investigator obtains data through intervention or interaction, or whose identifiable private information is used (U.S. Department of Health and Human Services, 2018). A rater's score can be obtained *from* a person without being *about* that person.

Canadian policy offers no such phrase, and the objection should be put at full strength. Canadian Institutes of Health Research et al. (2022) defines participants as individuals whose data or "responses to interventions, stimuli or questions by the researcher, are relevant to answering the research question(s)". A rater's judgment is a response to a question, and it is plainly relevant to the research question, so on its face the definition takes in the evaluator. The distinction cannot rest on the bare fact that a judgment was requested.

It need not. The Canadian framework draws the same line, but by focus of inquiry rather than

by an *about* clause. Its own interpretation says that research relying only on information that staff “normally provide as part of their work duties” needs no review, because “the information is the focus of the research, not the views of the staff member” (Panel on Research Ethics, 2022). An expert asked whether an item meets a stated criterion is supplying information of just that kind: the focus is the item, not the assessor.

The same interpretation marks the boundary the argument has to respect. If researchers also ask those people “to provide personal opinions outside the scope of their job roles”, review is required (Panel on Research Ethics, 2022). The evaluator role therefore holds only while the judgment is read as a within-competence assessment of the object. Once a judgment is taken as the person’s own reaction, idiolect, or attitude, the focus has moved to the person, and the judge is a participant (section 5).

This locates the point about method correctly. That “choice of methodology and/or intent or ability to publish findings” do not determine whether an activity is research (Canadian Institutes of Health Research et al., 2022) blocks the inference from systematicity to participant status, but it does not by itself separate participant from evaluator. The focus-of-inquiry test does that work. The rater scores the response; the response does not score the rater.

7 REPORTING AND ETHICS

The practical proposal is a reporting distinction that follows the role. Participant studies should report recruitment, eligibility, consent, relevant demographics, compensation, exclusion, and the statistical treatment of responses. Evaluation studies should report a different set (American Educational Research Association et al., 2014):

- the evaluator’s qualifications for the standard applied;
- the evaluation task and the criterion or diagnostic used;
- the materials and context available to the evaluator;
- whether the evaluator was blind to the hypothesis;
- independence from the author and the project;
- how disagreements were adjudicated;
- any conflict of interest or relationship to the work;
- whether the judgment concerns form, meaning, register, dialect, processing difficulty, or theoretical classification.

This routing does not weaken ethics review; it aims it. A study that targets persons, whether their processing, dialect, attitudes, or the social distribution of their judgments, belongs in participant review with the usual protections. A study that uses expert competence to assess linguistic materials belongs under the standards that govern measurement: calibration, reliability, independence, and conflict management (American Educational Research Association et al., 2014).

The institutional claim stays narrow. Local boards decide the scope of their own review, and no definition settles policy. The conceptual point, argued in section 6, is only that the word “judgment” should not trigger participant review on its own. What should trigger it is whether the design studies persons as persons.

8 CONCLUSION

Expert grammaticality judgments should be classified by evidential role. When the research target is a population of judgments, an effect on judges, or the social distribution of acceptability, judges

are participants. When the target is the linguistic status of materials and expert judgment is used to evaluate those materials, judges are evaluators.

The payoff is methodological hygiene. It lets grammaticality research be transparent about expert judgment without pretending that every expert consultation is a human-participant study. The right analogy is often not the survey respondent but the essay scorer: a fallible, reportable, correctable evaluator whose work affects measurement quality. Where the judgment is the object, review the judge; where the judgment is the instrument, calibrate the judgment.

ACKNOWLEDGEMENTS

The large language models GPT-5.5 and Claude Opus 4.8 served as drafting and editing aids throughout the preparation of this paper. I am responsible for all theoretical claims, arguments, errors, and interpretive choices.

REFERENCES

- American Educational Research Association, American Psychological Association & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*.
- Canadian Institutes of Health Research, Natural Sciences and Engineering Research Council of Canada & Social Sciences and Humanities Research Council of Canada. (2022). *Tri-council policy statement: Ethical conduct for research involving humans* (Policy statement) (TCPS 2 (2022)). Government of Canada. Retrieved June 18, 2026, from https://ethics.gc.ca/eng/policy-politique_tcps2-eptc2_2022.html
- Gibson, E., & Fedorenko, E. (2010). Weak quantitative standards in linguistics research. *Trends in Cognitive Sciences*, 14(6), 233–234. <https://doi.org/10.1016/j.tics.2010.03.005>
- Gibson, E., & Fedorenko, E. (2013). The need for quantitative methods in syntax and semantics research. *Language and Cognitive Processes*, 28(1-2), 88–124. <https://doi.org/10.1080/01690965.2010.515080>
- Gibson, E., Piantadosi, S. T., & Fedorenko, E. (2011). Using Mechanical Turk to obtain and analyze English acceptability judgments. *Language and Linguistics Compass*, 5(8), 509–524. <https://doi.org/10.1111/j.1749-818X.2011.00295.x>
- McCaffrey, D. F., Casabianca, J. M., Ricker-Pedley, K. L., Lawless, R. R., & Wendler, C. (2022). Best practices for constructed-response scoring [ETS Research Report RR-22-17]. *ETS Research Report Series*, 2022(1). <https://doi.org/10.1002/ets2.12358>
- Panel on Research Ethics. (2022). *TCPS 2 interpretations: Scope of REB review* [Interagency Advisory Panel on Research Ethics; “the information is the focus of the research, not the views of the staff member”]. Retrieved June 18, 2026, from https://ethics.gc.ca/eng/policy-politique_interpretations_scope-portee.html
- Schütze, C. T. (2016). *The empirical base of linguistics: Grammaticality judgments and linguistic methodology*. Language Science Press. <https://doi.org/10.17169/langsci.b89.100>
- Sprouse, J., Schütze, C. T., & Almeida, D. (2013). A comparison of informal and formal acceptability judgments using a random sample from *Linguistic Inquiry* 2001–2010. *Lingua*, 134, 219–248. <https://doi.org/10.1016/j.lingua.2013.07.002>
- U.S. Department of Health and Human Services. (2018). *Federal policy for the protection of human subjects (common rule)*, 45 CFR part 46 [Human-subject definition at 45 CFR 46.102(e)(1)]. Retrieved June 18, 2026, from <https://www.ecfr.gov/current/title-45/part-46>